

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
import pandas as pd
import numpy as np
import bz2
import re
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
from sklearn.utils import shuffle
from tqdm import tqdm
from keras.layers import *
from keras.models import Model
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
```

```
data1 = bz2.BZ2File('/content/gdrive/MyDrive/train.ft.txt.bz2')
```

```
data1 = data1.readlines()
```

```
data = [x.decode('utf-8') for x in data1]
```

```
data = data[:200000]
```

```
train_labels = [0 if x.split(' ')[0] == '__label__1' else 1 for x in data]
train_sentences = [x.split(' ', 1)[1][:-1].lower() for x in data]
```

```
for i in range(len(train_sentences)):
    if 'www.' in train_sentences[i] or 'http:' in train_sentences[i] or 'https:' in train_
        train_sentences[i] = re.sub(r"([ ]+(?<=\.[a-z]{3}))", "<url>", train_sentences[i])
```

```
data = pd.DataFrame(train_sentences, columns = ['review'])
data.head()
```

#### review

- |   | review  |
|---|---|
| 0 | stuning even for the non-gamer: this sound tra... |
| 1 | the best soundtrack ever to anything.: i'm rea... |
| 2 | amazing!: this soundtrack is my favorite music... |
| 3 | excellent soundtrack: i truly like this soundt... |
| 4 | remember, pull your jaw off the floor after he... |

```
data['score'] = train_labels  
data.head()
```

	review	score
0	stuning even for the non-gamer: this sound tra...	1
1	the best soundtrack ever to anything.: i'm rea...	1
2	amazing!: this soundtrack is my favorite music...	1
3	excellent soundtrack: i truly like this soundt...	1
4	remember, pull your jaw off the floor after he...	1

```
# De algemene info over de dataset  
data.describe()
```

	score
count	200000.000000
mean	0.505830
std	0.499967
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

```
# De dimensies van de dataset  
data.shape
```

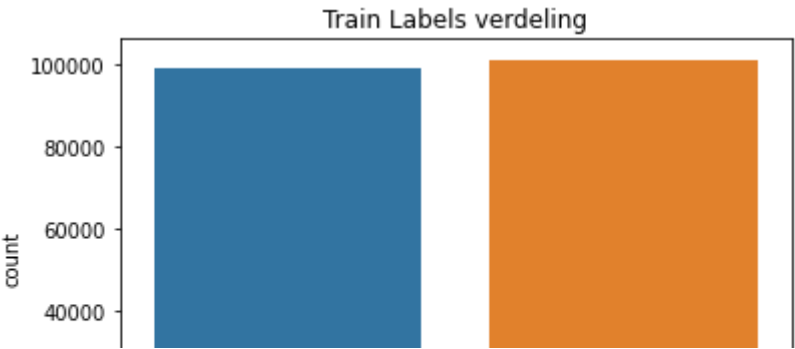
```
(200000, 2)
```

```
# Het aantal records  
data.count()
```

```
review    200000  
score     200000  
dtype: int64
```

```
sns.countplot(train_labels)  
plt.title('Train Labels verdeling')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass
FutureWarning
Text(0.5, 1.0, 'Train Labels verdeling')
```



```
data['word_count'] = data['review'].apply(lambda x : len(x.split()))
data['char_count'] = data['review'].apply(lambda x : len(x.replace(" ", "")))
data['word_density'] = data['word_count'] / (data['char_count'] + 1)
```

```
data.head()
```

	review	score	word_count	char_count	word_density
0	stuning even for the non-gamer: this sound tra...	1	80	347	0.229885
1	the best soundtrack ever to anything.: i'm rea...	1	97	413	0.234300
2	amazing!: this soundtrack is my favorite music...	1	129	632	0.203791
	excellent soundtracks i truly like this				

```
data.describe()
```


	score	word_count	char_count	word_density
count	200000.000000	200000.000000	200000.000000	200000.000000
mean	0.505830	80.092535	361.827595	0.222160
std	0.499967	43.251276	197.531975	0.021086
min	0.000000	10.000000	78.000000	0.035294
25%	0.000000	44.000000	194.000000	0.208556
50%	1.000000	72.000000	323.000000	0.222222
75%	1.000000	111.000000	500.000000	0.235669
max	1.000000	241.000000	873.000000	0.444444

```
fig, ax = plt.subplots(1, 3, figsize=(16, 6))
dp=sns.distplot(data['word_count'],ax=ax[0])
dp=sns.distplot(data['char_count'],ax=ax[1])
dp=sns.distplot(data['word_density'],ax=ax[2])
plt.show()
```

```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:
warnings.warn(msg, FutureWarning)

```



```

# Het verwijderen van de stopwoorden
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import string
nltk.download('punkt')

def verwijderStopwoorden(text):
    stopwoorden = set(stopwords.words('english'))
    tokens = word_tokenize(text.lower())
    result = [x for x in tokens if x not in stopwoorden]
    seperator = ' '
    return seperator.join(result)

data['review'] = data['review'].map(verwijderStopwoorden)

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.

```

```
data['review'].head()
```

```

0    stuning even non-gamer : sound track beautiful...
1    best soundtrack ever anything . : 'm reading l...
2    amazing ! : soundtrack favorite music time , h...

```

```

3    excellent soundtrack : truly like soundtrack e...
4    remember , pull jaw floor hearing : 've played...
Name: review, dtype: object

```

```
dataRF = data
```

```
dataRF = dataRF.drop(['word_count', 'char_count','word_density'], axis=1)
```

```
try:
```

```
    %tensorflow_version 2.x
```

```
except Exception:
```

```
    pass
```

```
import tensorflow as tf
```

```
from tensorflow import keras
```

```
print(tf.__version__)
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import sklearn as sk
```

```
import pandas as pd
```

```
seed = 2020
```

```
np.random.seed(seed)
```

```
import sklearn as sk
```

```
from sklearn.model_selection import train_test_split
```

```
from tensorflow.keras.datasets import mnist
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense, Dropout, Flatten, Embedding, Conv1D, MaxPoolin
```

```
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.constraints import max_norm
```

```
from tensorflow.keras.models import load_model
```

```
import nltk
```

## 2.4.1

```
def plot_history(history):
```

```
    plt.figure(figsize = (12,4))
```

```
    plt.subplot(1,2,1)
```

```
    plt.xlabel('Epoch')
```

```
    plt.ylabel('Nauwkeurigheid')
```

```
    plt.plot(history.epoch, np.array(history.history['accuracy']), 'g-',
             label='Train dataset nauwkeurigheid')
```

```
    plt.plot(history.epoch, np.array(history.history['val_accuracy']), 'r-',
             label = 'Validatie dataset nauwkeurigheid')
```

```
    plt.legend()
```

```
plt.subplot(1,2,2)
plt.xlabel('Epoch')
plt.ylabel('Verlies')
plt.plot(history.epoch, np.array(history.history['loss']), 'g-',
         label='Train dataset verlies')
plt.plot(history.epoch, np.array(history.history['val_loss']), 'r-',
         label = 'Validatie dataset verlies')
plt.legend()
```

```
# De lengte van de woorden in een nieuwe kolom
dataRF['numberOfWords'] = dataRF.review.str.split().apply(len)
dataRF.head()
```

	review	score	numberOfWords
0	stuning even non-gamer : sound track beautiful...	1	51
1	best soundtrack ever anything . : 'm reading l...	1	58
2	amazing ! : soundtrack favorite music time , h...	1	108
3	excellent soundtrack : truly like soundtrack e...	1	101
4	remember , pull jaw floor hearing : 've played...	1	68

```
dataRF['numberOfWords'].describe()
```

```
count    200000.000000
mean      53.914610
std       28.933542
min        7.000000
25%       30.000000
50%       48.000000
75%       74.000000
max      433.000000
Name: numberOfWords, dtype: float64
```

```
# Een training dataset opzetten
from sklearn.model_selection import train_test_split
X = dataRF.drop(['score', 'numberOfWords'], axis=1)
y = dataRF['score']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30)
```

```
X_train = np.asarray(X_train)
X_test = np.asarray(X_test)
```

```
num_classes = 2
#De methode utils.to_categorical zet vectors om in binaire matrices
#De scores (0 of 1) worden omgezet naar een binaire matrix. Het aantal klassen is
#twee omdat er twee opties zijn: positief of negatief
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```

from tensorflow.keras.layers.experimental.preprocessing import TextVectorization
#TextVectorization zet een lijst van strings om in een lijst van tokens
vectorizer = TextVectorization(max_tokens=20000, output_sequence_length=20)
text_ds = tf.data.Dataset.from_tensor_slices(X_train).batch(128)
vectorizer.adapt(text_ds)

voc = vectorizer.get_vocabulary()
word_index = dict(zip(voc, range(len(voc))))

#Glove staat voor Global Vector
glove_file = '/content/gdrive/My Drive/glove.6B.100d.txt'

#Glove-files bevatten woord vectors. De file die hier gebruikt wordt, bevat 400.000 vector
embeddings_index = {}
with open(glove_file) as f:
    for line in f:
        values = line.split()
        woord = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[woord] = coefs

num_tokens = len(voc) + 2
embedding_dim = 100
missed_words = []

# Een embedding matrix aanmaken
embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
    else:
        missed_words.append(word)

num_classes = 2
#Het model initialiseren
def initial_model():
    model = Sequential()

    model.add(Embedding(num_tokens, embedding_dim, embeddings_initializer=keras.initialize
    model.add(Conv1D(16,activation='relu',kernel_size=3))
    model.add(MaxPooling1D(3))
    model.add(Dropout(0.2))

    model.add(Conv1D(16,activation='relu',kernel_size=3))
    model.add(Dropout(0.2))

    model.add(Conv1D(16,activation='relu',kernel_size=3))
    model.add(GlobalMaxPooling1D())
    model.add(Dense(16, activation='relu', kernel_initializer='he_uniform'))

```

```
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

#Categorical Crossentropy berekent het cross-entropie verlies tussen de labels en de voor  
#Dit is een optimalisator die het Adam-algoritme implementeert.

```
model.compile(loss='categorical_crossentropy',
              optimizer=tf.keras.optimizers.Adam(learning_rate = 0.0001),
              metrics=['accuracy'])
return model
```

```
#Het model verwacht een array, dus dit wordt hier omgezet
X_train_final = vectorizer(np.array([s for s in X_train])).numpy()
X_test_final = vectorizer(np.array([s for s in X_test])).numpy()
```

```
y_train_final = np.array(y_train)
y_test_final = np.array(y_test)
```

```
model_1 = initial_model()
model_1.summary()
batch_size = 128
epochs = 30
```

```
history_1 = model_1.fit(X_train_final, y_train_final,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(X_test_final, y_test_final)
                        )
```

```
Total params: 2,006,890
Trainable params: 6,690
Non-trainable params: 2,000,200
```

```
Epoch 1/30
1094/1094 [=====] - 14s 12ms/step - loss: 0.7144 - accuracy: 0.1250
Epoch 2/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.6293 - accuracy: 0.1875
Epoch 3/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.5611 - accuracy: 0.2500
Epoch 4/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.5131 - accuracy: 0.3125
Epoch 5/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4926 - accuracy: 0.3750
Epoch 6/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4749 - accuracy: 0.4375
Epoch 7/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4654 - accuracy: 0.5000
Epoch 8/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4570 - accuracy: 0.5625
Epoch 9/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4544 - accuracy: 0.6250
Epoch 10/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4468 - accuracy: 0.6875
Epoch 11/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4456 - accuracy: 0.7500
Epoch 12/30
```



```

Epoch 12/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4400 - accuracy: 0.7600
Epoch 13/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4356 - accuracy: 0.7656
Epoch 14/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4315 - accuracy: 0.7715
Epoch 15/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4279 - accuracy: 0.7779
Epoch 16/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4252 - accuracy: 0.7852
Epoch 17/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4221 - accuracy: 0.7921
Epoch 18/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4219 - accuracy: 0.7919
Epoch 19/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4211 - accuracy: 0.7911
Epoch 20/30
1094/1094 [=====] - 13s 12ms/step - loss: 0.4142 - accuracy: 0.7942
Epoch 21/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4173 - accuracy: 0.7973
Epoch 22/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4142 - accuracy: 0.7942
Epoch 23/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4096 - accuracy: 0.7996
Epoch 24/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4096 - accuracy: 0.7996
Epoch 25/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4104 - accuracy: 0.7984
Epoch 26/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4078 - accuracy: 0.7958
Epoch 27/30
1094/1094 [=====] - 12s 11ms/step - loss: 0.4076 - accuracy: 0.7956
Epoch 28/30

```

#De resultaten visualiseren

```

[train_loss, train_accuracy] = model_1.evaluate(X_train_final, y_train_final, verbose=0)
print("Nauwkeurigheid training dataset:{:7.2f}".format(train_accuracy))
[val_loss, val_accuracy] = model_1.evaluate(X_test_final, y_test_final, verbose=0)
print("Nauwkeurigheid test dataset:{:7.2f}".format(val_accuracy))
plot_history(history_1)

```

↳ Nauwkeurigheid training dataset: 0.83  
 Nauwkeurigheid test dataset: 0.82

