

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
import pandas as pd
import numpy as np
import bz2
import re
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
from sklearn.utils import shuffle
from tqdm import tqdm
from keras.layers import *
from keras.models import Model
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
```

```
data = pd.read_csv('/content/gdrive/MyDrive/IMDB Dataset.csv', encoding="ISO-8859-1")
```

```
data['review'][2]
```

'I thought this was a wonderful way to spend time on a too hot summer weekend, sitting in the air conditioned theater and watching a light-hearted comedy. The plot is simplistic, but the dialogue is witty and the characters are likable (even the well bread suspected serial killer). While some

```
# De dimensies van de dataset
```

```
data.shape
```

```
(50000, 2)
```

```
# Het aantal records
```

```
data.count()
```

```
review      50000
sentiment    50000
dtype: int64
```

```
# Het verwijderen van de stopwoorden
```

```
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import string
nltk.download('punkt')
```

Opgeslagen.



```
stopwords = set(stopwords.words('english'))
tokens = word_tokenize(text.lower())
```

```
tokens = word_tokenize(text.lower())
result = [x for x in tokens if x not in stopwoorden and not x.startswith('@')]
separator = ' '
return separator.join(result)
```

```
data['review'] = data['review'].map(verwijderStopwoorden)
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

```
data['sentiment'][data['sentiment']=='negative'] = 0
data['sentiment'][data['sentiment']=='positive'] = 1
```

```
def clean_html(text):
    clean = re.compile('<.*?>')
    return re.sub(clean, '',text)
```

```
data['review'] = data['review'].apply(clean_html)
data.head()
```

|   | review  | sentiment |
|---|---|-----------|
| 0 | one reviewers mentioned watching 1 oz episode ... | 1         |
| 1 | wonderful little production . filming techni...   | 1         |
| 2 | thought wonderful way spend time hot summer we... | 1         |
| 3 | basically 's family little boy ( jake ) thinks... | 0         |
| 4 | better motto is `` love time money `` visual      | 1         |

```
data['review'][7]
```

```
'show amazing , fresh & innovative idea 70 's first aired . first 7 8 yea
rs brilliant , things dropped . 1990 , show really funny anymore , 's con
tinued decline complete waste time today. 's truly disgraceful far show
```

```
data.describe()
```

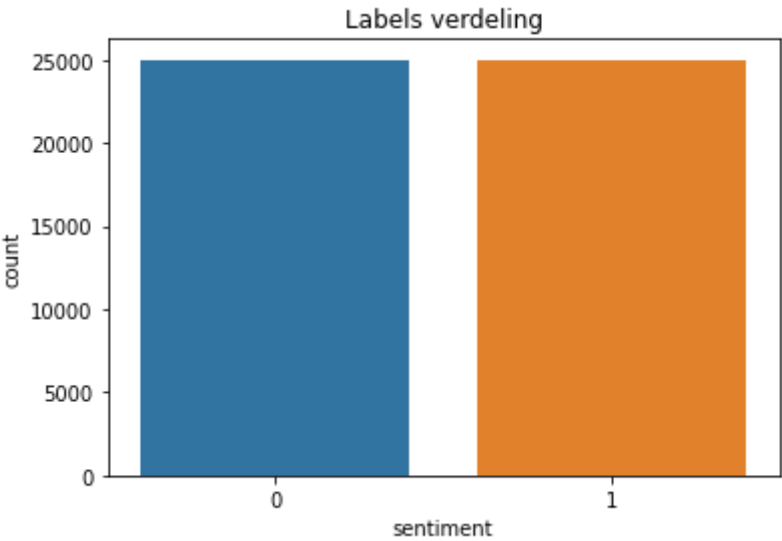
|        | review  | sentiment |
|--------|---|-----------|
| count  | 50000   | 50000     |
| unique | 49578   | 2         |
| top    | loved today 's show !!! variety solely cooki... | 1         |
| freq   | 5   | 25000     |

Opgeslagen.

Verken

```
sns.countplot(data['sentiment'])
plt.title('Labels verdeling')

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning
FutureWarning
Text(0.5, 1.0, 'Labels verdeling')
```



```
data['word_count'] = data['review'].apply(lambda x : len(x.split()))
data['char_count'] = data['review'].apply(lambda x : len(x.replace(" ","")))
data['word_density'] = data['word_count'] / (data['char_count'] + 1)
```

```
data.head()
```

|   | review   | sentiment | word_count | char_count | word_density |
|---|--|-----------|------------|------------|--------------|
| 0 | one reviewers mentioned watching 1 oz episode ...                        | 1         | 213        | 1029       | 0.206796     |
| 1 | wonderful little production . filming techni...<br>thought wonderful way | 1         | 105        | 601        | 0.174419     |

```
data.describe()
```


Opgeslagen.

×

|       | word_count   | char_count   | word_density |
|-------|--------------|--------------|--------------|
| count | 50000.000000 | 50000.000000 | 50000.000000 |

```
fig, ax = plt.subplots(1, 3, figsize=(16, 6))
dp=sns.distplot(data['word_count'],ax=ax[0])
dp=sns.distplot(data['char_count'],ax=ax[1])
dp=sns.distplot(data['word_density'],ax=ax[2])
plt.show()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: FutureWarning
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: FutureWarning
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: FutureWarning
warnings.warn(msg, FutureWarning)
```



```
dataRF = data
```

```
dataRF = dataRF.drop(['word_count', 'char_count', 'word_density'], axis=1)
```

```
try:
```

```
    %tensorflow_version 2.x
except Exception:
    pass
```

```
import tensorflow as tf
from tensorflow import keras
```

```
print(tf.__version__)
```

Opgeslagen.

```
import matplotlib.pyplot as plt
```

```

import sklearn as sk
import pandas as pd

seed = 2020
np.random.seed(seed)

import sklearn as sk
from sklearn.model_selection import train_test_split

from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Embedding, Conv1D, MaxPoolin
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.constraints import max_norm
from tensorflow.keras.models import load_model

import nltk

```

VERBORGEN UITVOER WEERGEVEN

```

def plot_history(history):
    plt.figure(figsize = (12,4))
    plt.subplot(1,2,1)

    plt.xlabel('Epoch')
    plt.ylabel('Nauwkeurigheid')
    plt.plot(history.epoch, np.array(history.history['accuracy']), 'g-',
             label='Train dataset nauwkeurigheid')
    plt.plot(history.epoch, np.array(history.history['val_accuracy']), 'r-',
             label = 'Validatie dataset nauwkeurigheid')
    plt.legend()

    plt.subplot(1,2,2)
    plt.xlabel('Epoch')
    plt.ylabel('Verlies')
    plt.plot(history.epoch, np.array(history.history['loss']), 'g-',
             label='Train dataset verlies')
    plt.plot(history.epoch, np.array(history.history['val_loss']), 'r-',
             label = 'Validatie dataset verlies')
    plt.legend()

# De lengte van de woorden in een nieuwe kolom
dataRF['numberOfWords'] = dataRF.review.str.split().apply(len)
dataRF.head()

```

Opgeslagen.



|   | review  | sentiment | numberOfWords |
|---|---|-----------|---------------|
| 0 | one reviewers mentioned watching 1 oz episode ... | 1         | 213           |

```
dataRF['numberOfWords'].describe()
```

```
count    50000.000000
mean      152.845840
std       115.158279
min        5.000000
25%       82.000000
50%      114.000000
75%      186.000000
max      1705.000000
Name: numberOfWords, dtype: float64
```

```
# Een training dataset opzetten
```

```
from sklearn.model_selection import train_test_split
```

```
X = dataRF.drop(['sentiment','numberOfWords'],axis=1)
```

```
y = dataRF['sentiment']
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.30)
```

```
X_train = np.asarray(X_train)
```

```
X_test = np.asarray(X_test)
```

```
num_classes = 2
```

```
#De methode utils.to_categorical zet vectors om in binaire matrices
```

```
#De scores (0 of 1) worden omgezet naar een binaire matrix. Het aantal klassen is  
#twee omdat er twee opties zijn: positief of negatief
```

```
y_train = keras.utils.to_categorical(y_train, num_classes)
```

```
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
from tensorflow.keras.layers.experimental.preprocessing import TextVectorization
```

```
#TextVectorization zet een lijst van strings om in een lijst van tokens
```

```
vectorizer = TextVectorization(max_tokens=20000, output_sequence_length=20)
```

```
text_ds = tf.data.Dataset.from_tensor_slices(X_train).batch(128)
```

```
vectorizer.adapt(text_ds)
```

```
voc = vectorizer.get_vocabulary()
```

```
word_index = dict(zip(voc, range(len(voc))))
```

```
#Glove staat voor Global Vector
```

```
glove_file = '/content/gdrive/My Drive/glove.6B.100d.txt'
```

```
#Glove-files bevatten woord vectors. De file die hier gebruikt wordt, bevat 400.000 vector
```

```
embeddings_index = {}
```

```
with open(glove_file) as f:
```

```
    for line in f:
```

```
        values = line.split()
```

Opgeslagen.



dtype='float32')

s

```

num_tokens = len(voc) + 2
embedding_dim = 100
missed_words = []

# Een embedding matrix aanmaken
embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
    else:
        missed_words.append(word)

#Het model initialiseren
def initial_model3():
    model = Sequential()

    model.add(Embedding(num_tokens, embedding_dim,
                        embeddings_initializer=keras.initializers.Constant(embedding_matrix)))
    model.add(LSTM(2))

    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model

def initial_model4():
    model = Sequential()

    model.add(Embedding(num_tokens, embedding_dim,
                        embeddings_initializer=keras.initializers.Constant(embedding_matrix)))
    model.add(Dropout(0.20))
    model.add(Conv1D(16, activation='relu', kernel_size=2))
    model.add(GlobalMaxPooling1D())
    model.add(Dropout(0.20))
    model.add(Dense(num_classes, activation='sigmoid'))

    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model

def initial_model2():
    model = Sequential()

```

Opgeslagen.



```

embedding_dim,
initializer=keras.initializers.Constant(embedding_matrix))
model.add(Bidirectional(LSTM(16, dropout=0.2, recurrent_dropout=0.2)))

```

```

model.add(Dense(32, activation='relu'))
model.add(Dropout(0.50))
model.add(Dense(2, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

return model

```

```

#Het model verwacht een array, dus dit wordt hier omgezet
X_train_final = vectorizer(np.array([s for s in X_train])).numpy()
X_test_final = vectorizer(np.array([s for s in X_test])).numpy()

y_train_final = np.array(y_train)
y_test_final = np.array(y_test)

model_1 = initial_model4()
model_1.summary()
batch_size = 128
epochs = 30

history_1 = model_1.fit(X_train_final, y_train_final,
                       batch_size=batch_size,
                       epochs=epochs,
                       verbose=1,
                       validation_data=(X_test_final, y_test_final)
                       )

```

Model: "sequential\_1"

| Layer (type)                    | Output Shape      | Param # |
|---------------------------------|-------------------|---------|
| embedding_1 (Embedding)         | (None, None, 100) | 2000200 |
| dropout_1 (Dropout)             | (None, None, 100) | 0       |
| conv1d_1 (Conv1D)               | (None, None, 16)  | 3216    |
| global_max_pooling1d_1 (Glob    | (None, 16)        | 0       |
| dropout_2 (Dropout)             | (None, 16)        | 0       |
| dense_1 (Dense)                 | (None, 2)         | 34      |
| Total params: 2,003,450         |                   |         |
| Trainable params: 3,250         |                   |         |
| Non-trainable params: 2,000,200 |                   |         |

Epoch 1/30

274/274 [=====] - 3s 10ms/step - loss: 0.7905 - accuracy:

Epoch 2/30

[=====] - 2s 9ms/step - loss: 0.6281 - accuracy: (

[=====] - 2s 9ms/step - loss: 0.5982 - accuracy: (

Opgeslagen.





```

Epoch 4/30
274/274 [=====] - 2s 9ms/step - loss: 0.5825 - accuracy: 0.70
Epoch 5/30
274/274 [=====] - 2s 9ms/step - loss: 0.5784 - accuracy: 0.71
Epoch 6/30
274/274 [=====] - 2s 9ms/step - loss: 0.5726 - accuracy: 0.72
Epoch 7/30
274/274 [=====] - 2s 9ms/step - loss: 0.5683 - accuracy: 0.73
Epoch 8/30
274/274 [=====] - 2s 9ms/step - loss: 0.5659 - accuracy: 0.73
Epoch 9/30
274/274 [=====] - 2s 9ms/step - loss: 0.5578 - accuracy: 0.74
Epoch 10/30
274/274 [=====] - 2s 9ms/step - loss: 0.5586 - accuracy: 0.74
Epoch 11/30
274/274 [=====] - 2s 9ms/step - loss: 0.5562 - accuracy: 0.74
Epoch 12/30
274/274 [=====] - 2s 9ms/step - loss: 0.5508 - accuracy: 0.74
Epoch 13/30
274/274 [=====] - 2s 9ms/step - loss: 0.5517 - accuracy: 0.74
Epoch 14/30
274/274 [=====] - 2s 9ms/step - loss: 0.5469 - accuracy: 0.74
Epoch 15/30
274/274 [=====] - 2s 9ms/step - loss: 0.5510 - accuracy: 0.74
Epoch 16/30
274/274 [=====] - 2s 9ms/step - loss: 0.5436 - accuracy: 0.74
Epoch 17/30
274/274 [=====] - 2s 9ms/step - loss: 0.5461 - accuracy: 0.74
Epoch 18/30
274/274 [=====] - 2s 9ms/step - loss: 0.5404 - accuracy: 0.74
Epoch 19/30
274/274 [=====] - 2s 9ms/step - loss: 0.5461 - accuracy: 0.74

```

#De resultaten visualiseren

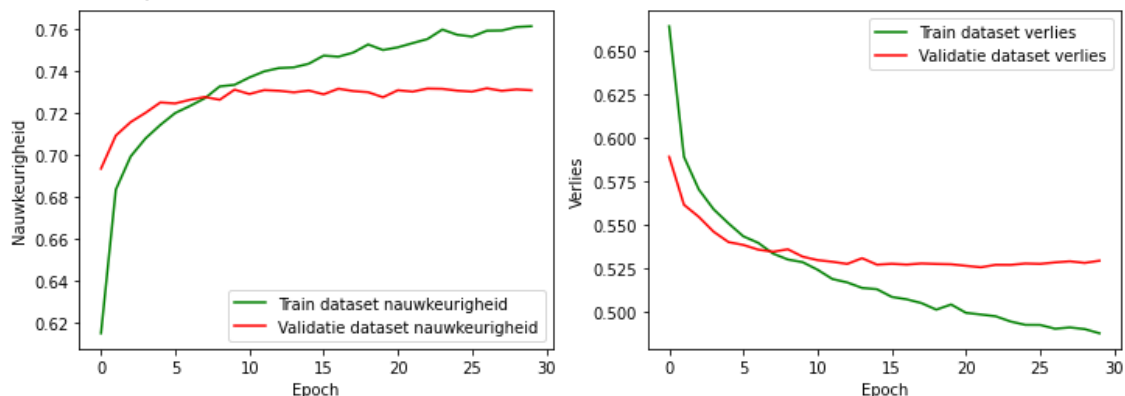
```

[train_loss, train_accuracy] = model_1.evaluate(X_train_final, y_train_final, verbose=0)
print("Nauwkeurigheid training dataset:{:7.2f}".format(train_accuracy))
[val_loss, val_accuracy] = model_1.evaluate(X_test_final, y_test_final, verbose=0)
print("Nauwkeurigheid test dataset:{:7.2f}".format(val_accuracy))
plot_history(history_1)

```

Nauwkeurigheid training dataset: 0.79

Nauwkeurigheid test dataset: 0.73



Opgeslagen.

